

Gaussian Naive Bayes for Classification

DATA 2060 Final Project

Professor Andras Zsom

Axle Zhao, Jiaqi Zhang, Ghirish Thaenraj

Overview - Gaussian Naive Bayes

- Probabilistic *generative* classifier based on Bayes' Theorem
- Assumes conditional independence of features
- Continuous features modeled with Gaussian distributions
- Very fast, interpretable, and surprisingly strong baseline
- Common in: medical diagnosis, spam filtering, real-time systems

Bayes' Theorem for Classification

Bayes' Theorem:

$$P(y | x) = \frac{P(x | y)P(y)}{P(x)}$$

The set up of our learning problem:

$$P(y = k | x) = \frac{P(y = k)P(x | y = k)}{P(x)} = \frac{P(y = k) \prod_{i=1}^d P(x_i | y = k)}{P(x)}$$

where

$$\begin{aligned} P(x) &= \sum_k P(x \cap (y = k)) = \sum_k P(y = k) \prod_{i=1}^d P(x_i | y = k) \\ P(x | y = k) &= \prod_{i=1}^d P(x_i | y = k) \end{aligned}$$

because we assume features are independent given the class label

Since $P(x)$ is constant:

$$\hat{y} = \arg \max_k \left[\log P(y = k) + \sum_{j=1}^d \log P(x_j | y = k) \right]$$

Note: We work in log-space for numerical stability. (To prevent underflow)

Gaussian Likelihood Model

$P(y = k)$ is the fraction of the input data that has the label of k

For each feature x_j under class k :

$$P(x_j|y = k) = \frac{1}{\sqrt{2\pi\sigma_{jk}^2}} \exp\left(-\frac{(x_j - \mu_{jk})^2}{2\sigma_{jk}^2}\right)$$

μ_{jk} = mean of x_j over samples with $y = k$

σ_{jk}^2 = variance of x_j over samples with $y = k$.

Numerical Stability & Smoothing

Why smoothing?

- Zero variance → division by zero in Gaussian PDF
- Solve by adding smoothing value ϵ

Sklearn-style smoothing:

$$\epsilon = 10^{-9} \cdot \max(\text{Var}(X))$$

Key Numerical Techniques

- Maximum Likelihood Estimation (MLE) for parameters
- Log-space computation to avoid underflow
- Variance smoothing to prevent division-by-zero
- Gaussian log-PDF for stable likelihood calculations
- Log-Sum-Exp normalization for probability outputs
- Fully vectorized operations for efficiency

Training Algorithm (Pseudo-Code)

Input:

```
X (n × d matrix of features)  
y (n labels)
```

For each **class k**:

```
X_k ← all samples in X with label k  
prior[k] ← |X_k| / n  
mean[k] ← average of each feature in X_k  
var[k] ← variance of each feature in X_k + epsilon # smoothing
```

Output:

```
priors, means, variances
```

Prediction Algorithm (Pseudo-Code)

Input:

x ($1 \times d$ feature vector), priors, means, variances

For each **class** k :

$\text{log_prior}[k] \leftarrow \log(\text{prior}[k])$

$\text{log_likelihood}[k] \leftarrow 0$

For each feature j :

$\text{log_likelihood}[k] \leftarrow \text{log_likelihood}[k]$

$+ \log \text{Gaussian_PDF}(x_j ; \text{mean}[k,j], \text{var}[k,j])$

$\text{log_posterior}[k] \leftarrow \text{log_prior}[k] + \text{log_likelihood}[k]$

Return the **class with** the maximum log_posterior

Probability Normalization

To compute $P(y=k | x)$:

Shift all log_posterior values by subtracting their maximum

Compute $\exp(\text{log_posterior_shifted}[k])$ for each class

Normalize by dividing by the sum of exponentials

Unit Testing Overview


We implemented granular tests for:

- Input reshaping: $1D \rightarrow 2D$ (Typo, we somehow included it by accident.)
- Zero-variance features
- Single-class datasets
- Loss function correctness
- Log-likelihood computation
- Probability normalization


Outcome: all tests pass and validate correct numerical behavior.

Accuracy Comparison (Iris Dataset)

- Dataset: Iris (multiclass, 4 continuous features)
- Train/test split: 70/30, fixed random state
- Our model accuracy: 97.78%
- Sklearn accuracy: 97.78%
- All predicted class labels match exactly




Confusion Matrix — GNB Model Performance



Probability Comparison

Probabilities match Sklearn to within:

$$\text{max difference} = 5.55 \times 10^{-16}$$

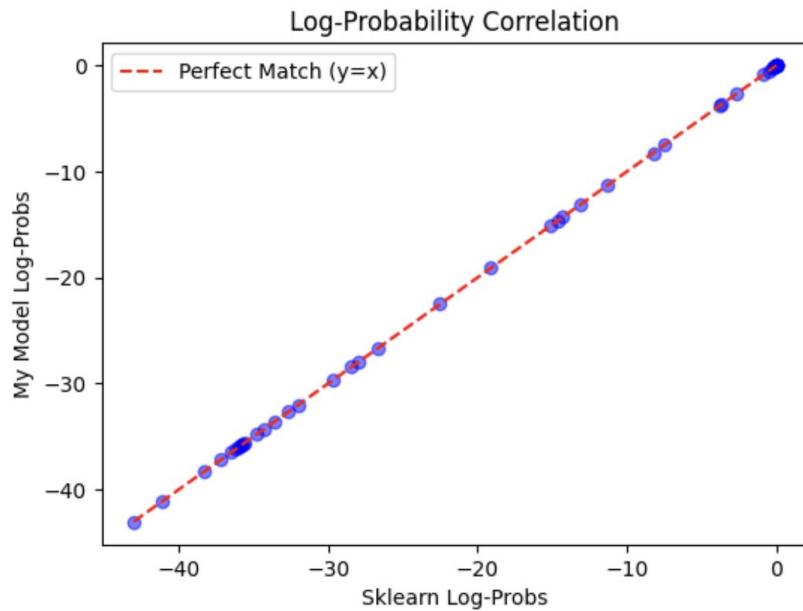


Log Probability Comparison

Log probabilities also match to machine precision

Confirms correct implementation of:


- Gaussian log-PDF
- Log-sum-exp normalization
- Class prior incorporation



GNB Decision Boundaries (2D Projection)


Generating 2D Decision Boundary (Sepal Length vs Width)

Gaussian Naive Bayes Decision Boundaries (2D)



Epsilon Smoothing Verification

- Our computed epsilon = Sklearn's epsilon
- This was crucial for exact probability reproduction
- Confirms correct replication of Sklearn internals



Key Takeaways

- GNB is simple but powerful
- Independence assumption → extremely efficient
- Generative models provide full probability distributions
- Numerical stability (log-space + smoothing) is essential
- Our implementation reproduces Sklearn exactly

Challenges We Encountered

- Understanding Sklearn's dynamic variance smoothing
- Handling zero-variance features
- Normalizing log probabilities correctly
- Ensuring vectorization without losing numerical stability
- Designing comprehensive unit tests

Summary

We accomplished:

- Full scratch implementation of Gaussian Naive Bayes
- Correct mathematical modeling
- Robust unit tests covering edge cases
- Exact match with Sklearn in:
 - Predictions
 - Probabilities
 - Log-probabilities
 - Smoothing calculation
- Clear visual verification through our graphs

Thank You